

# Podman Community Call:

- AI-Lab podman-desktop extension
- ai-lab-recipes

June 4, 2024

Sally O'Malley

Principal Software Engineer

Emerging Technologies, Office of CTO

@sallyom

---

# AI journey for an **Application Developer**

Announcement:

---

# Come to DevConf.US in Boston Aug 14-16, it's FREE!!

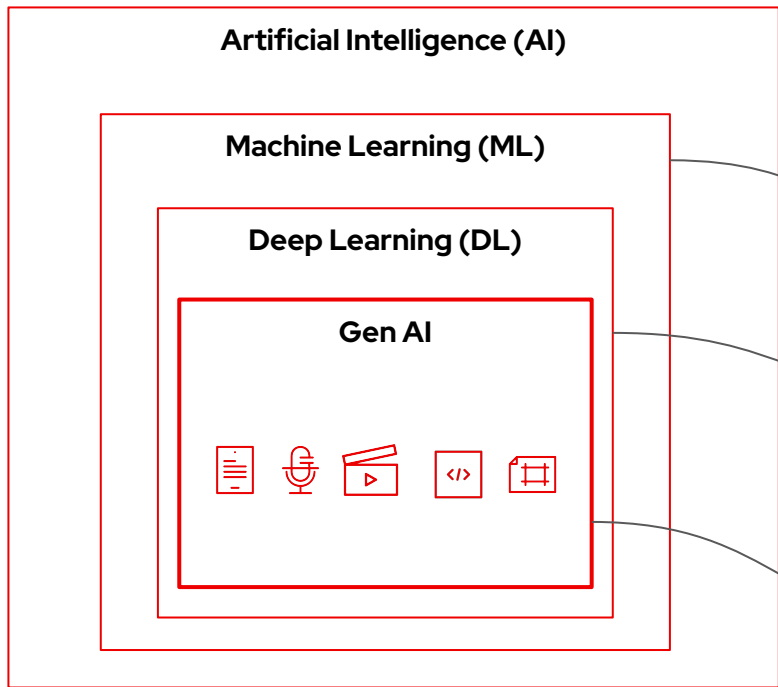
- [DevConf.US](https://devconf.us)
- Containerization Guild will be there, Wed  
5:30 PM

---

# Watch Cedric's presentation!

## Generative AI Development from Local to Production

# What is Generative AI?



## Artificial Intelligence

A multidisciplinary field of Computer Science that aims to create systems capable of emulating and surpassing human-level intelligence

## Machine Learning

A subfield for training computers to learn and make decisions from data, without explicit programming

## Deep Learning

A subset from ML involved neural networks enabling complex pattern recognition and analysis from data

## Gen AI Learning

Subfield focusing on generating new content such as text, images, code and more


# What is Generative AI?

## Generative AI

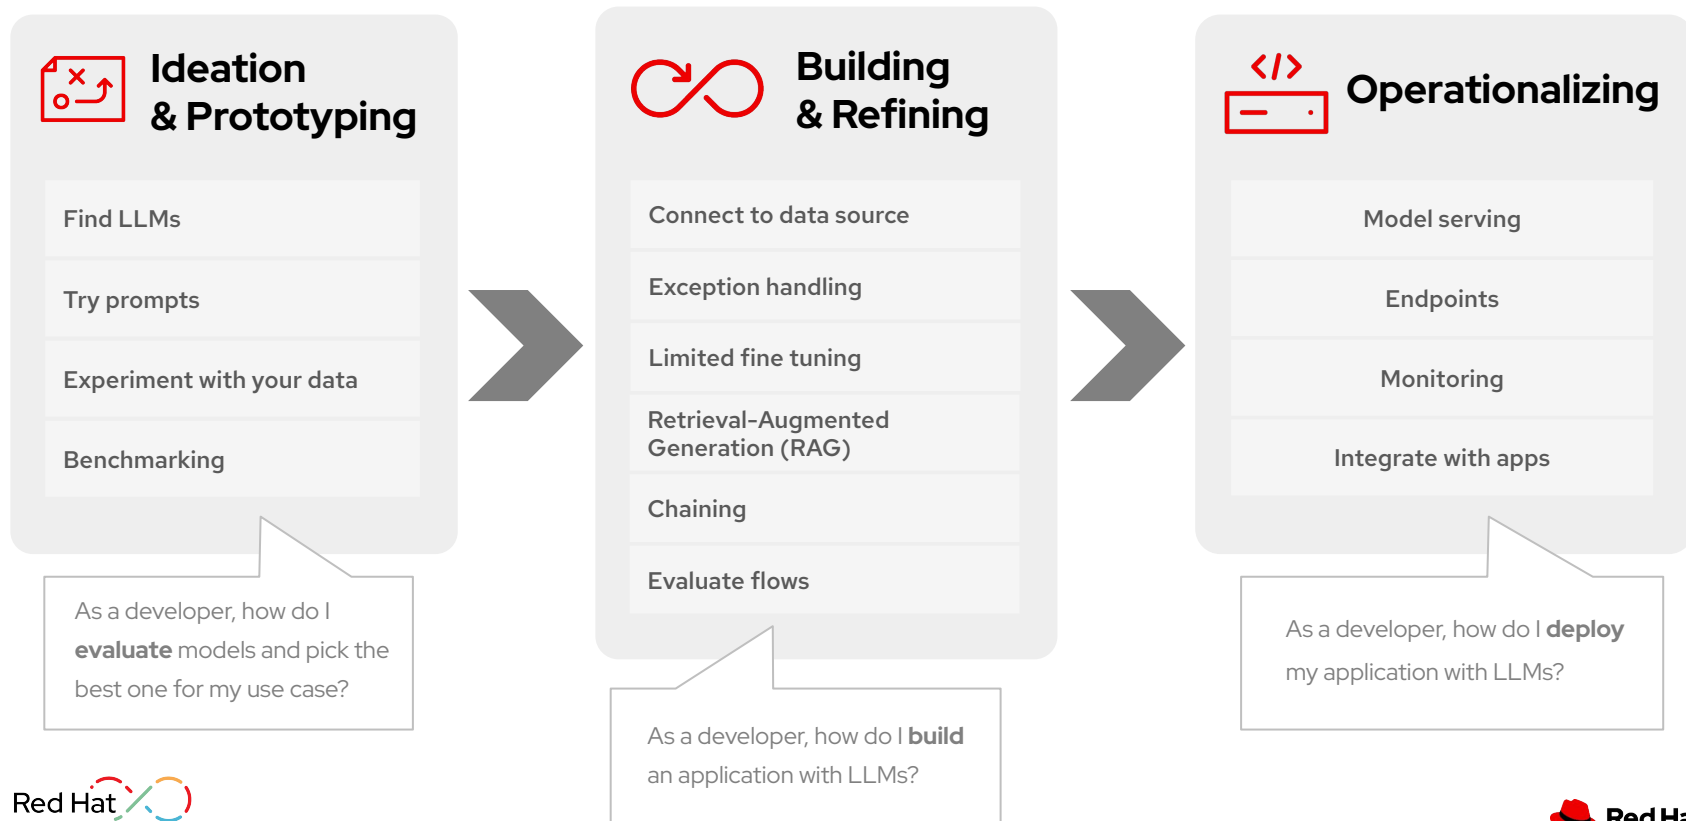
*Generative AI utilizes AI models to create text, images, videos, code, or data in response to prompts*

## LLMs

*Large Language Models are a specialized type of AI, trained on massive datasets to generate outputs in natural, conversational language*




# Adopting Generative AI for an **Application Developer**



---

# How do I get started?

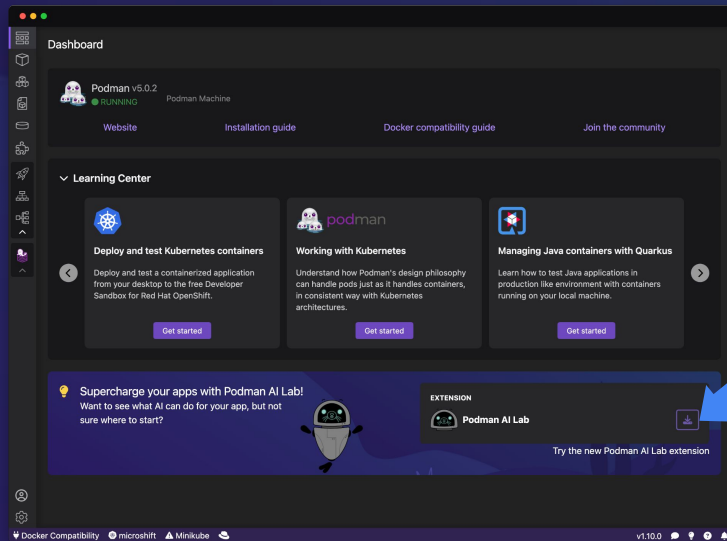




# Podman AI

Lab

Run AI MLs locally and build AI applications



[github.com/containers/podman](https://github.com/containers/podman)  
[-desktop-extension-ai-lab](#)

Supported platforms:



# Introducing: Podman AI Lab

## Your developer environment for working with GenAI

### Discover GenAI

- Get inspired by AI use cases
- Learn how to integrate AI in an optimal way
- Experiment with different compatible Models

### Model Catalog

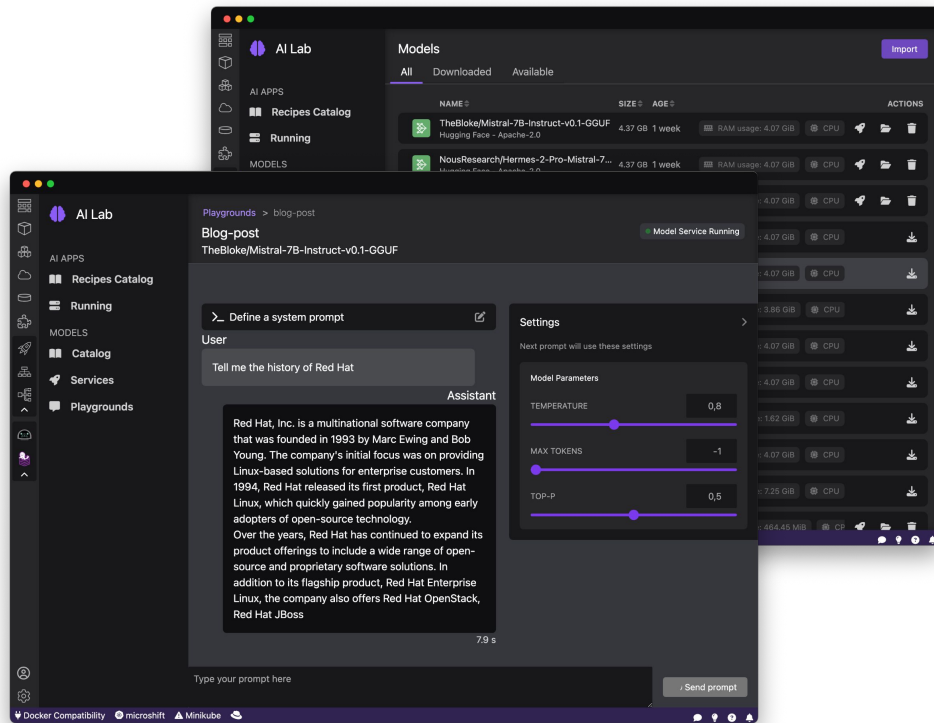
- Leverage a curated list of open source large language models available out of the box
- Import your own models

### Run Models Locally

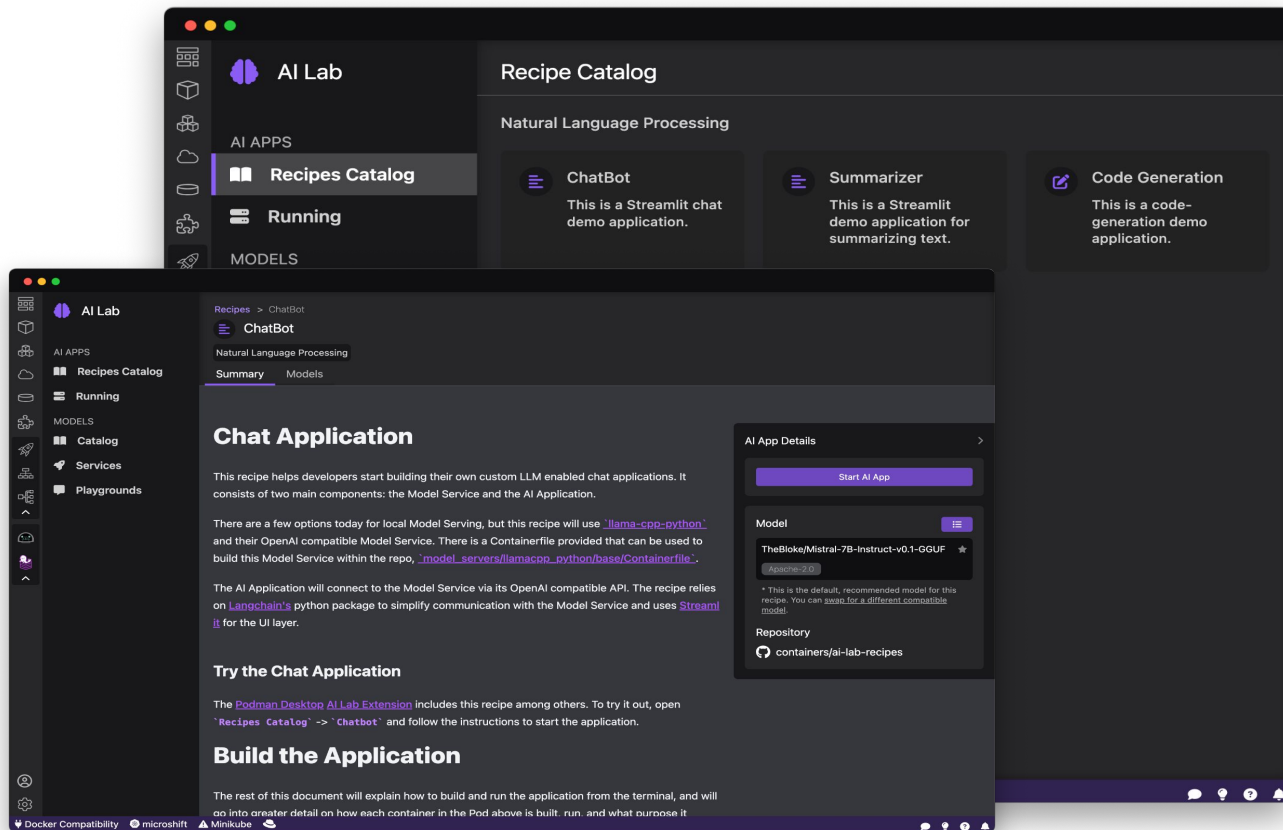
- Run models with an inference server running in UBI image
- Get OpenAI compatible API
- Use code snippets

### Playground Environment

- Experiment with models and prompts
- Configure settings and system prompts
- Test and validate prompt workflows before using in your application



# Recipes



# Model Catalog

The screenshot displays the AI Lab Model Catalog interface. The left sidebar contains navigation options: AI APPS, Recipes Catalog, Running, MODELS, Catalog (selected), Services, and Playgrounds. The main area shows a list of models under the 'All' tab, with columns for NAME, SIZE, AGE, and ACTIONS. The models listed are:

| NAME  | SIZE    | AGE    | ACTIONS                           |
|---|---------|--------|-----------------------------------|
| TheBloke/Mistral-7B-Instruct-v0.1-GGUF<br>Hugging Face - Apache-2.0 | 4.37 GB | 1 week | RAM usage: 4.07 GiB, CPU, [Icons] |
| NousResearch/Hermes-2-Pro-Mistral-7...                              | 4.37 GB | 1 week | RAM usage: 4.07 GiB, CPU, [Icons] |
| ibm/merlinite-7b-GGUF<br>Hugging Face                               |         |        | [Icons]                           |
| TheBloke/Mi...  |         |        | [Icons]                           |
| TheBloke/Mi...  |         |        | [Icons]                           |
| froggeric/Ce...   |         |        | [Icons]                           |
| TheBloke/op...  |         |        | [Icons]                           |
| TheBloke/Mi...  |         |        | [Icons]                           |
| MazyarPana...   |         |        | [Icons]                           |
| llmware/drag...   |         |        | [Icons]                           |
| MazyarPana...   |         |        | [Icons]                           |
| ggerganov/v...  |         |        | [Icons]                           |

The right panel shows the 'Service Details' for the 'ibm/merlinite-7b-GGUF' model. It includes a 'Container' section with a unique ID, a 'Models' section showing the model name and license, a 'Server' section with the URL and inference type, and a 'Client code' section with a Python script for interacting with the service.

```
python
Python LangChain

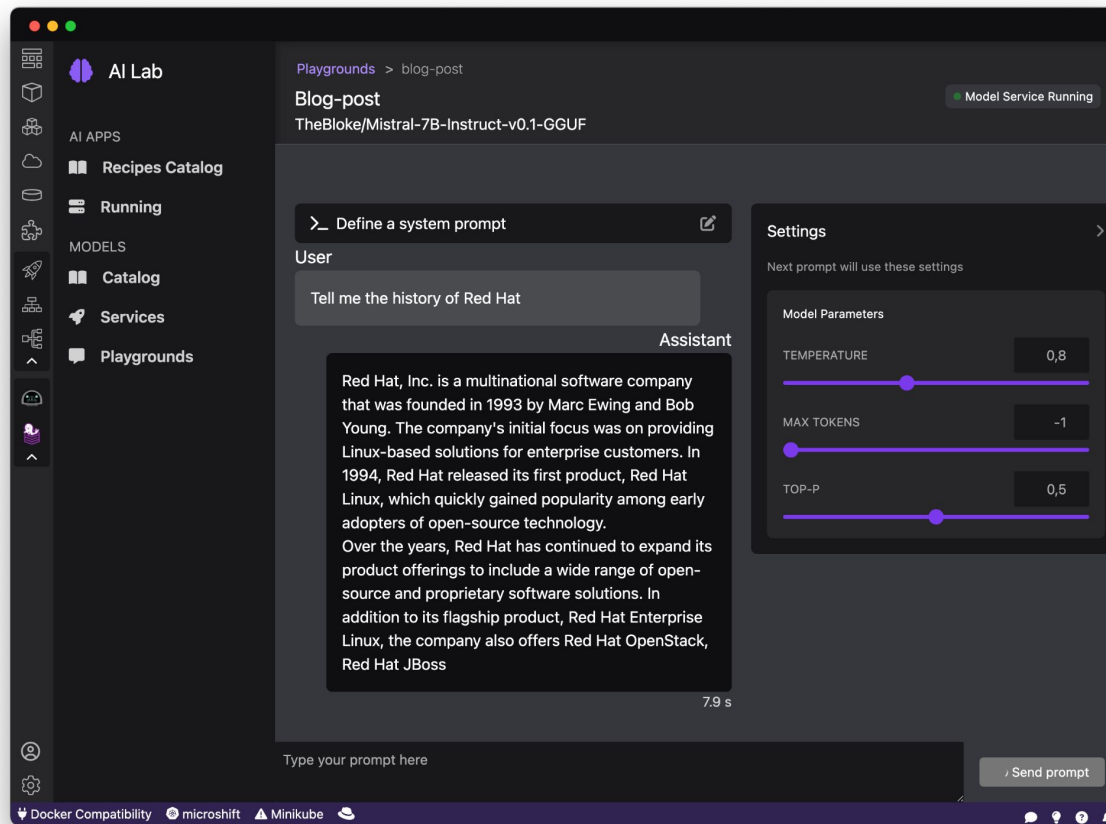
pip
=====
pip install langchain langchain-openai

AiService.py
=====
from langchain_openai import OpenAI
from langchain.chains import LLMChain
from langchain_core.prompts import ChatPromptTemplate

model_service = "http://localhost:61067/v1/"

llm = OpenAI(base_url=model_service,
```

# Playground Environments



---

# Demo Time!

---

**Tell me MORE about the recipes!**

**ai-lab-recipes**

<https://github.com/containers/ai-lab-recipes>

---

# Demo Time!



---

# Blogs!

- [podman-desktop ai-lab extension](#)
- [ai-lab-recipes](#)

---

## Other resources!

- [Andrej Karpathy interview](#)
- [Ollama](#)
- [LocalAI](#)

# Get started with Podman & Podman Desktop!

Get involved with Podman and adopt a seal (figuratively)!



## On the Web

Learn how to install Podman

[podman.io](https://podman.io)

[podman-desktop.io](https://podman-desktop.io)




## Github

Contribute and report issues

[github.com/containers/podman](https://github.com/containers/podman)

[github.com/containers/podman-desktop](https://github.com/containers/podman-desktop)



## Chat

Get connected with the community

[podman.io/community](https://podman.io/community)

# Thank you

## Join the DevNation

Red Hat Developer serves the builders. The problem solvers who create careers with code. Let's keep in touch!

- Join Red Hat Developer at [developers.redhat.com/register](https://developers.redhat.com/register)
- Follow us on any of our social channels
- Visit [dn.dev/upcoming](https://dn.dev/upcoming) for a schedule of our upcoming events

## Red Hat Developer

Build here. Go anywhere.



[linkedin.com/company/red-hat](https://linkedin.com/company/red-hat)



[youtube.com/user/RedHatVideos](https://youtube.com/user/RedHatVideos)



[facebook.com/redhatinc](https://facebook.com/redhatinc)



[twitter.com/RedHat](https://twitter.com/RedHat)